



PLURALSIGHT

M&T Bank Data Academy

Week 5



Tarek Atwan
Instructor, Pluralsight

Proprietary and confidential

What is a Database

A database is a structured collection of data that is organized and stored in a computer system. It is designed to efficiently manage, store, and retrieve large amounts of data for various applications and users.

- **Data organization:** Databases use a structured format to organize data into tables, rows, and columns, allowing for efficient storage and retrieval.
- **Data integrity:** Databases enforce data integrity by implementing rules and constraints that ensure the accuracy and consistency of the stored data.
- **Data security:** Databases provide mechanisms for controlling access to the data, ensuring that only authorized users can view, modify, or delete the stored information.
- **Data concurrency:** Databases support multiple users accessing and modifying the data simultaneously, while maintaining data consistency and integrity.
- **Data recovery:** Databases implement backup and recovery mechanisms to protect against data loss and ensure the availability of the stored information.

What is a Data Warehouse

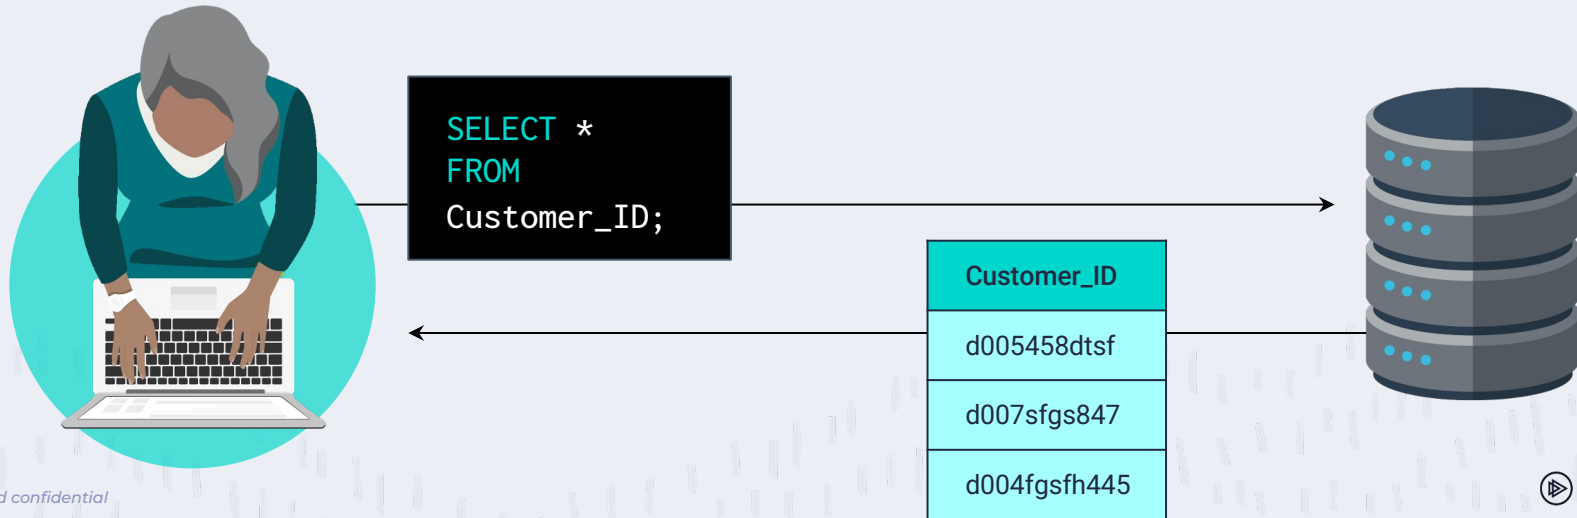
A data warehouse (DW) is a type of database that is specifically designed for analytical rather than transactional work. It collects and aggregates data from various sources, such as operational databases, external data feeds, and more, to provide a consolidated and integrated view of the organization's data.

- **Structured and filtered data:** Data warehouses store structured, filtered, and pre-processed data that is optimized for analytical queries and reporting.
- **Historical data:** Data warehouses often include historical data, allowing users to analyze trends and patterns over time.
- **Separation of operational and analytical workloads:** Data warehouses are separate from operational databases, ensuring that analytical queries do not impact the performance of transactional systems.
- **Support for complex queries and reporting:** Data warehouses are optimized for complex analytical queries and reporting, allowing users to gain insights from large volumes of data.

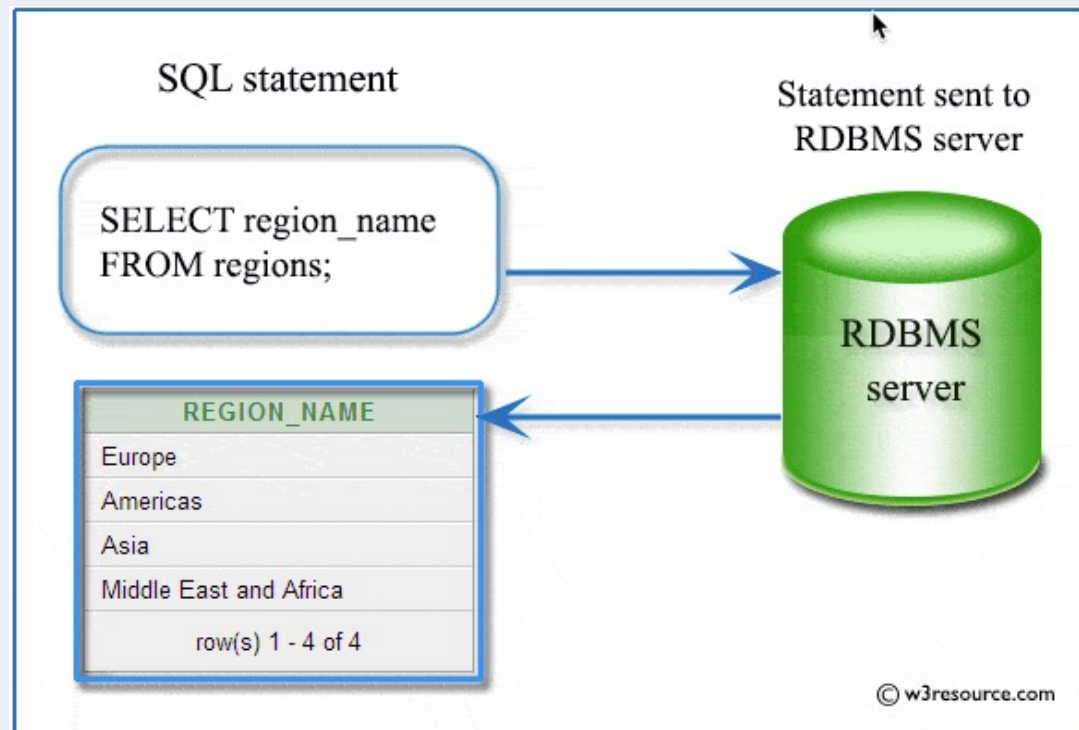
What is SQL

SQL (often pronounced "sequel") stands for Structured Query Language.

It is a powerful tool that enables programmers to create, populate, manipulate, and access databases. It also provides an easy method for dealing with server-side storage.



What is SQL



What is SQL

Data using SQL is stored in tables on the server, much like spreadsheets you would create in Microsoft Excel.

This makes the data easy to visualize and search.



Customer_ID	Date_ID
d005458dtsf	6/26/2019
d007sfgs847	8/3/2018
d004fgsfh445	12/3/2018

Order_ID	Customer_ID	Date_ID
10001	d005458dtsf	6/26/2019
10002	d007sfgs847	8/3/2018
10003	d004fgsfh445	12/3/2018

Using CRUD

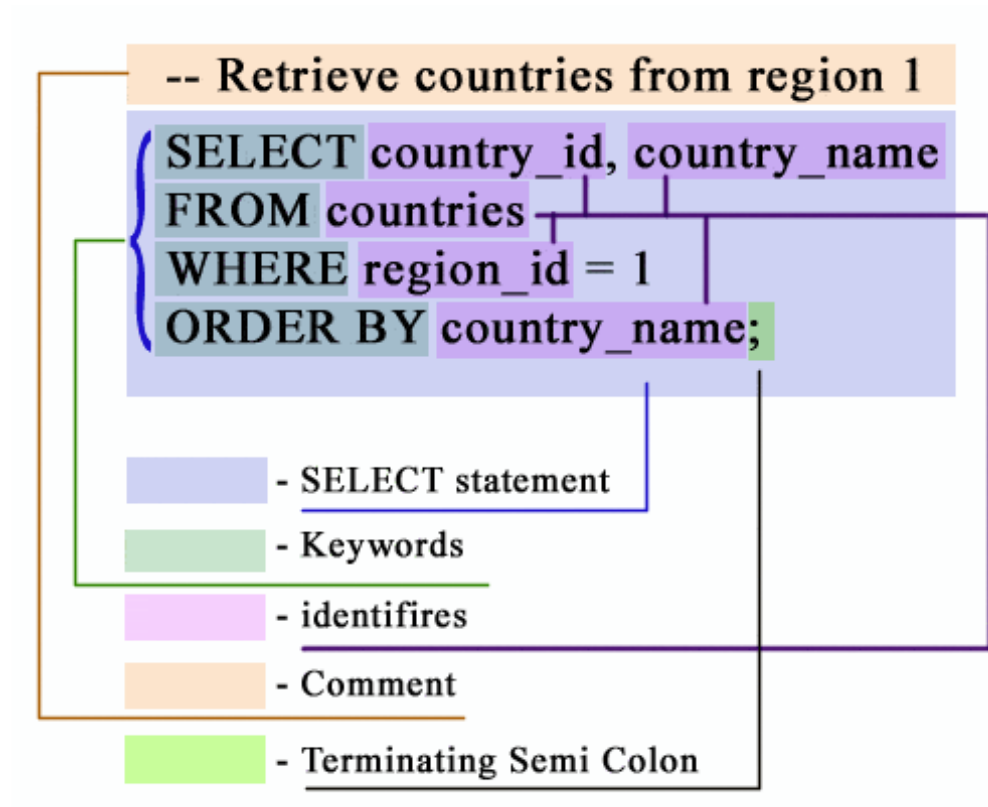
Create Read Update Delete is a set of operations used with persistent storage.

Create	Create data in a table with the INSERT statement.
Read	Read data by using SELECT .
Update	Updated a table's data by using UPDATE .
Delete	Deleted data via DELETE .

ANATOMY OF SQL QUERY (READ)

- SELECT** Used to specify the **columns** or expression you want to retrieve from the database
- FROM** Used to specify the **table** from which you want to retrieve the data
- WHERE** An optional clause used to filter the rows (data) returned
- ORDER BY** An optional clause used to sort the rows returned by the query based on one or more columns.

SQL Query Elements



Query Example

The **SELECT** clause can specify more than one column.

```
SELECT pet_type, pet_name  
FROM people  
WHERE pet_type = 'dog'  
AND pet_age < 5;
```

Query Example

Data is filtered by using additional clauses such as **WHERE** and **AND**.

```
SELECT pet_type, pet_name  
FROM people  
WHERE pet_type = 'dog'  
AND pet_age < 5;
```

Query Example

The **WHERE** clause will extract only the data that meets the condition specified. **AND** adds a second condition to the original clause, further refining the query.

```
SELECT pet_type, pet_name  
FROM people  
WHERE pet_type = 'dog'  
AND pet_age < 5;
```

Query Example

Note that unlike in Python where comparisons are done with a double equals (`==`) sign, in SQL only a single equal sign is used.

```
SELECT pet_type, pet_name  
FROM people  
WHERE pet_type = 'dog'  
AND pet_age < 5;
```

Wildcard: % and _

Use wildcards to substitute zero, one, or multiple characters in a string. The keyword **LIKE** indicates the use of a wildcard.

```
SELECT *  
FROM actor  
WHERE last_name LIKE 'Will%';
```

Wildcard: % and _

The **%** will substitute zero, one, or multiple characters in a query.
In this example, all of the following are matches: Will, Willa, and Willows.

```
SELECT *  
FROM actor  
WHERE last_name LIKE 'Will%';
```


Wildcard: % and _

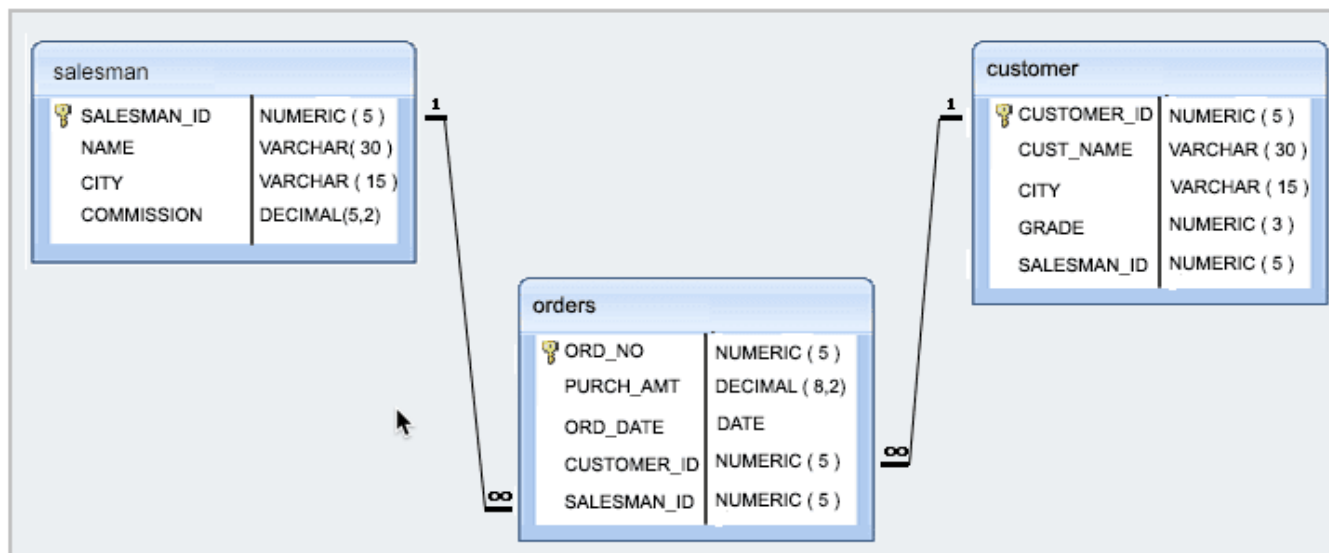
The `_` will substitute only **one** character in a query.

`_an` returns all actors whose first name contains three letters, the second and third of which are `an`.

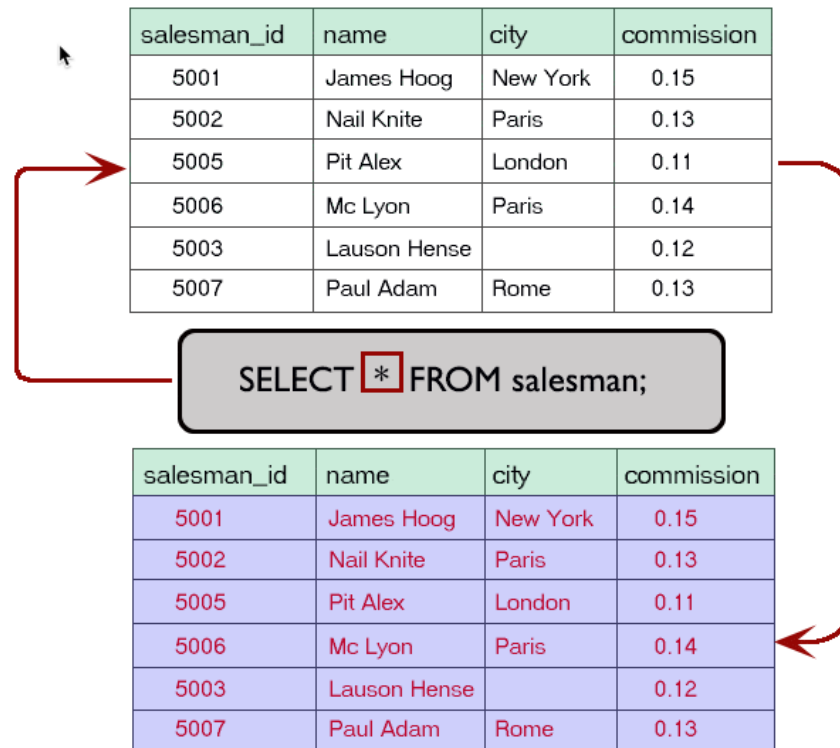
```
SELECT *  
FROM actor  
WHERE first_name LIKE '_an';
```

SELECT *

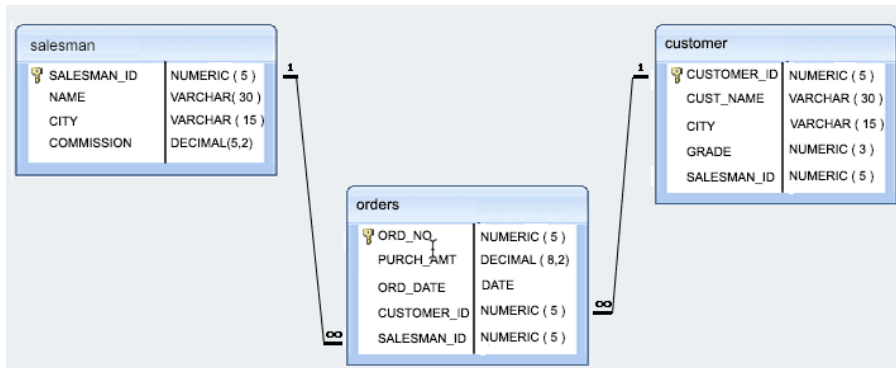
Structure of inventory database :



SELECT *



SELECT COLUMNS IN DIFFERENT ORDER

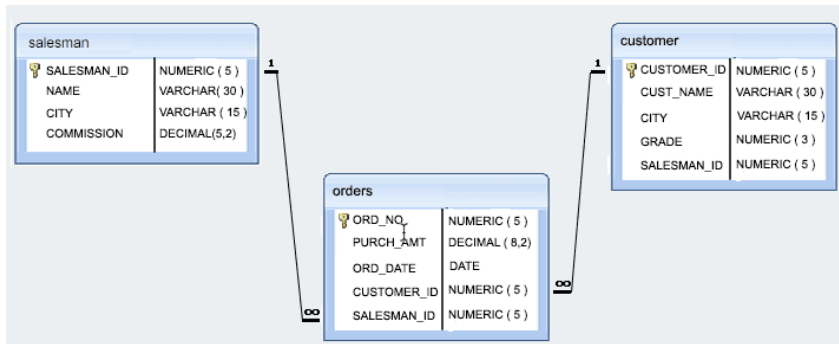


ord_no	purch_amt	ord_date	customer_id	salesman_id
70001	150.50	2012-10-05	3005	5002
70009	270.65	2012-09-10	3001	5005
70002	65.26	2012-10-05	3002	5001
70004	110.50	2012-08-17	3009	5003
70007	948.50	2012-09-10	3005	5002
70005	2400.60	2012-07-27	3007	5001
70008	5760.00	2012-09-10	3002	5001
70010	1983.43	2012-10-10	3004	5006
70003	2480.40	2012-10-10	3009	5003
70012	250.45	2012-06-27	3008	5002
70011	75.29	2012-08-17	3003	5007
70013	3045.60	2012-04-25	3002	5001

```
SELECT ord_date,salesman_id,ord_no,purch_amt
FROM orders;
```

ord_date	salesman_id	ord_no	purch_amt
2012-10-05	5002	70001	150.50
2012-09-10	5005	70009	270.65
2012-10-05	5001	70002	65.26
2012-08-17	5003	70004	110.50
2012-09-10	5002	70007	948.50
2012-07-27	5001	70005	2400.60
2012-09-10	5001	70008	5760.00
2012-10-10	5006	70010	1983.43
2012-10-10	5003	70003	2480.40
2012-06-27	5002	70012	250.45

UNIQUE VALUES




ord_no	purch_amt	ord_date	customer_id	salesman_id
70001	150.50	2012-10-05	3005	5002
70009	270.65	2012-09-10	3001	5005
70002	65.26	2012-10-05	3002	5001
70004	110.50	2012-08-17	3009	5003
70007	948.50	2012-09-10	3005	5002
70005	2400.60	2012-07-27	3007	5001
70008	5760.00	2012-09-10	3002	5001
70010	1983.43	2012-10-10	3004	5006
70003	2480.40	2012-10-10	3009	5003
70012	250.45	2012-06-27	3008	5002
70011	75.29	2012-08-17	3003	5007
70013	3045.60	2012-04-25	3002	5001

SELECT DISTINCT salesman_id
FROM orders;

salesman_id
5006
5002
5001
5005
5003
5007

SPECIFY A CONDITION

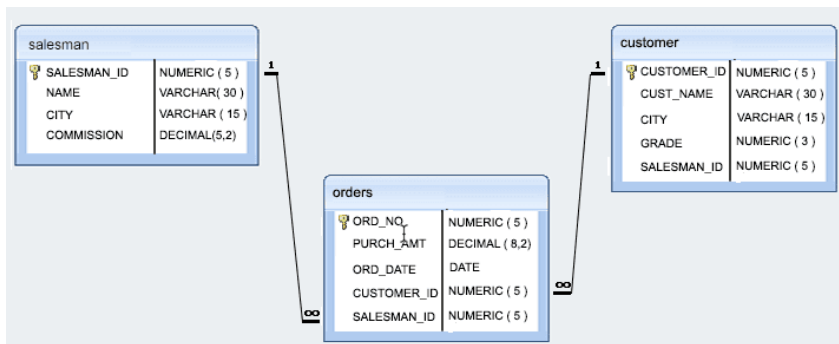


customer_id	cust_name	city	grade	salesman_id
3002	Nick Rimando	New York	100	5001
3005	Graham Zusi	California	200	5002
3004	Fabian Johnson	Paris	300	5006
3007	Brad Davis	New York	200	5001
3009	Geoff Cameron	Berlin	100	5003
3008	Julian Green	London	300	5002
3003	Jozy Altidore	Moscow	200	5007
3001	Brad Guzan	London		5005

```
SELECT * FROM customer
WHERE grade =200;
```

customer_id	cust_name	city	grade	salesman_id
3002	Nick Rimando	New York	100	5001
3005	Graham Zusi	California	200	5002
3004	Fabian Johnson	Paris	300	5006
3007	Brad Davis	New York	200	5001
3009	Geoff Cameron	Berlin	100	5003
3008	Julian Green	London	300	5002
3003	Jozy Altidore	Moscow	200	5007
3001	Brad Guzan	London		5005

COMBINING DATA




```

SELECT salesman.name AS "Salesman",
customer.cust_name, customer.city
FROM salesman, customer
WHERE salesman.city=customer.city;
  
```

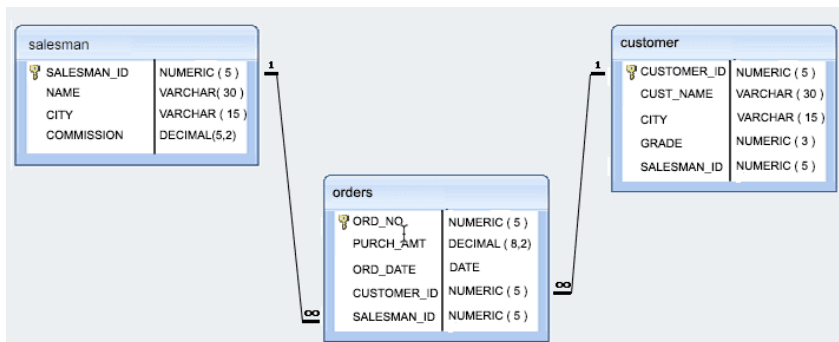
salesman_id	name	city	commission
5001	James Hoog	New York	0.15
5002	Nail Knite	Paris	0.13
5005	Pit Alex	London	0.11
5006	Mc Lyon	Paris	0.14
5003	Lauson Hense	London	0.12
5007	Paul Adam	Rome	0.13

customer_id	cust_name	city	grade	salesman_id
3002	Nick Rimando	New York	100	5001
3005	Graham Zusi	California	200	5002
3004	Fabian Johnson	Paris	300	5006
3007	Brad Davis	New York	200	5001
3009	Geoff Cameron	Berlin	100	5003
3008	Julian Green	London	300	5002
3001	Brad Guzan	London		5005
3003	Jozy Altidore	Moscow	200	5007



Salesman	cust_name	city
James Hoog	Brad Davis	New York
James Hoog	Nick Rimando	New York
Nail Knite	Fabian Johnson	Paris
Pit Alex	Brad Guzan	London
Pit Alex	Julian Green	London
Mc Lyon	Fabian Johnson	Paris

COMBINING DATA



```

SELECT a.ord_no,a.purch_amt,
b.cust_name,b.city
FROM orders a,customer b
WHERE a.customer_id=b.customer_id
AND a.purch_amt BETWEEN 500 AND 2000;
  
```

orders

ord_no	purch_amt	ord_date	customer_id	salesman_id
70001	150.50	2012-10-05	3005	5002
70009	270.65	2012-09-10	3001	5005
70002	65.26	2012-10-05	3002	5001
70004	110.50	2012-08-17	3009	5003
70007	948.50	2012-09-10	3005	5002
70005	2400.60	2012-07-27	3007	5001
70008	5760.00	2012-09-10	3002	5001
70010	1983.43	2012-10-10	3004	5006
70003	2480.40	2012-10-10	3009	5003
70012	250.45	2012-06-27	3008	5002
70011	75.29	2012-08-17	3003	5007
70013	3045.60	2012-04-25	3002	5001

customer

customer_id	cust_name	city	grade	salesman_id
3002	Nick Rimando	New York	100	5001
3005	Graham Zusi	California	200	5002
3004	Fabian Johnson	Paris	300	5006
3007	Brad Davis	New York	200	5001
3009	Geoff Cameron	Berlin	100	5003
3008	Julian Green	London	300	5002
3001	Brad Guzan	London	300	5005
3003	Jozy Altidore	Moscow	200	5007

orders

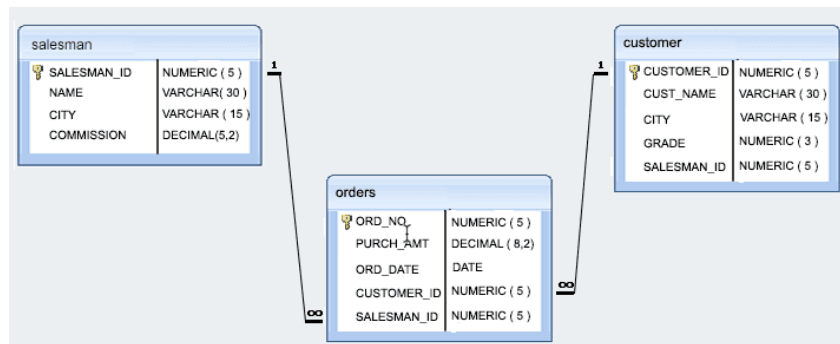
ord_no	purch_amt	ord_date	customer_id
70001	150.50	2012-10-05	3005
70009	270.65	2012-09-10	3001
70002	65.26	2012-10-05	3002
70004	110.50	2012-08-17	3009
70007	948.50	2012-09-10	3005
70005	2400.60	2012-07-27	3007
70008	5760.00	2012-09-10	3002
70010	1983.43	2012-10-10	3004

customer

customer_id	cust_name	city	grade	salesman_id
3002	Nick Rimando	New York	100	5001
3005	Graham Zusi	California	200	5002
3004	Fabian Johnson	Paris	300	5006
3007	Brad Davis	New York	200	5001
3009	Geoff Cameron	Berlin	100	5003
3008	Julian Green	London	300	5002
3001	Brad Guzan	London	300	5005
3003	Jozy Altidore	Moscow	200	5007

ord_no	purch_amt	cust_name	city
70007	948.50	Graham Zusi	California
70010	1983.43	Fabian Johnson	Paris

COMBINING DATA




```
SELECT a.cust_name AS "Customer Name",
a.city,b.name AS "Salesman", b.commission
FROM customer a
INNER JOIN salesman b
ON a.salesman_id=b.salesman_id;
```

customer				salesman			
customer_id	cust_name	city	grade	salesman_id	salesman_id	name	commission
3002	Nick Rimando	New York	100	5001	5001	James Hoog	0.15
3005	Graham Zusi	California	200	5002	5002	Nail Knite	0.13
3004	Fabian Johnson	Paris	300	5006	5006	Mc Lyon	0.14
3007	Brad Davis	New York	200	5001	5003	Lauson Hense	0.12
3009	Geoff Cameron	Berlin	100	5003	5005	Pit Alex	0.11
3008	Julian Green	London	300	5002	5007	Paul Adam	0.13
3001	Brad Guzan	London	5005	5005			
3003	Jozy Altidore	Moscow	200	5007			

customer			salesman_id	salesman		
cust_name	city	..	salesman_id	salesman_id	name	commission
Nick Rimando	New York	..	5001	5001	James Hoog	0.15
Brad Davis	New York	..	5001			
Graham Zusi	California	..	5002	5002	Nail Knite	0.13
Julian Green	London	..	5002			
Fabian Johnson	Paris	..	5006	5006	Mc Lyon	0.14
Geoff Cameron	Berlin	..	5003	5003	Lauson Hense	0.12
Brad Guzan	London	..	5005	5005	Pit Alex	0.11
Jozy Altidore	Moscow	..	5007	5007	Paul Adam	0.13

Customer Name	city	Salesman	commission
Nick Rimando	New York	James Hoog	0.15
Graham Zusi	California	Nail Knite	0.13
Fabian Johnson	Paris	Mc Lyon	0.14
Brad Davis	New York	James Hoog	0.15
Geoff Cameron	Berlin	Lauson Hense	0.12
Julian Green	London	Nail Knite	0.13
Brad Guzan	London	Pit Alex	0.11
Jozy Altidore	Moscow	Paul Adam	0.13

COMBINING DATA



```
SELECT a.cust_name AS "Customer Name",
a.city, b.name AS "Salesman", b.commission
FROM customer a
INNER JOIN salesman b
ON a.salesman_id=b.salesman_id
WHERE b.commission>.12;
```

customer

customer_id	cust_name	city	grade	salesman_id
3002	Nick Rimando	New York	100	5001
3005	Graham Zusi	California	200	5002
3004	Fabian Johnson	Paris	300	5006
3007	Brad Davis	New York	200	5001
3009	Geoff Cameron	Berlin	100	5003
3008	Julian Green	London	300	5002
3001	Brad Guzan	London	200	5005
3003	Jozy Altidore	Moscow	200	5007

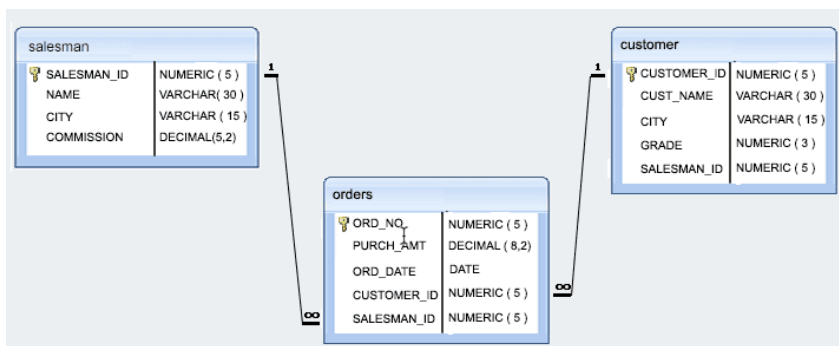
salesman

salesman_id	name	city	commission
5001	James Hoog	New York	0.15
5002	Nail Knite	Paris	0.13
5005	Pit Alex	London	0.11
5006	Mc Lyon	Paris	0.14
5003	Lauson Hense	Paris	0.12
5007	Paul Adam	Rome	0.13

customer				salesman		
cust_name	city	..	salesman_id	salesman_id	name	commission
Nick Rimando	New York	..	5001	5001	James Hoog	0.15
Brad Davis	New York	..	5001	5001	James Hoog	0.15
Graham Zusi	California	..	5002	5002	Nail Knite	0.13
Julian Green	London	..	5002	5002	Nail Knite	0.13
Fabian Johnson	Paris	..	5006	5006	Mc Lyon	0.14
Geoff Cameron	Berlin	..	5003	5003	Lauson Hense	0.12
Brad Guzan	London	..	5005	5005	Pit Alex	0.11
Jozy Altidore	Moscow	..	5007	5007	Paul Adam	0.13

Customer Name	city	Salesman	commission
Nick Rimando	New York	James Hoog	0.15
Graham Zusi	California	Nail Knite	0.13
Fabian Johnson	Paris	Mc Lyon	0.14
Brad Davis	New York	James Hoog	0.15
Julian Green	London	Nail Knite	0.13
Jozy Altidore	Moscow	Paul Adam	0.13

Find the total purchase amount for all orders




ord_no	purch_amt	ord_date	customer_id	salesman_id
70001	150.50	2012-10-05	3005	5002
70009	270.65	2012-09-10	3001	5005
70002	65.26	2012-10-05	3002	5001
70004	110.50	2012-08-17	3009	5003
70007	948.50	2012-09-10	3005	5002
70005	2400.60	2012-07-27	3007	5001
70008	5760.00	2012-09-10	3002	5001
70010	1983.43	2012-10-10	3004	5006
70003	2480.40	2012-10-10	3009	5003
70012	250.45	2012-06-27	3008	5002
70011	75.29	2012-08-17	3003	5007
70013	3045.60	2012-04-25	3002	5001

```
SELECT SUM(purch_amt)
FROM orders;
```

purch_amt
150.50
270.65
65.26
110.50
948.50
2400.60
5760.00
1983.43
2480.40
250.45
75.29
3045.60

Sum : 17541.18

Highest purchase amount ordered by the each customer



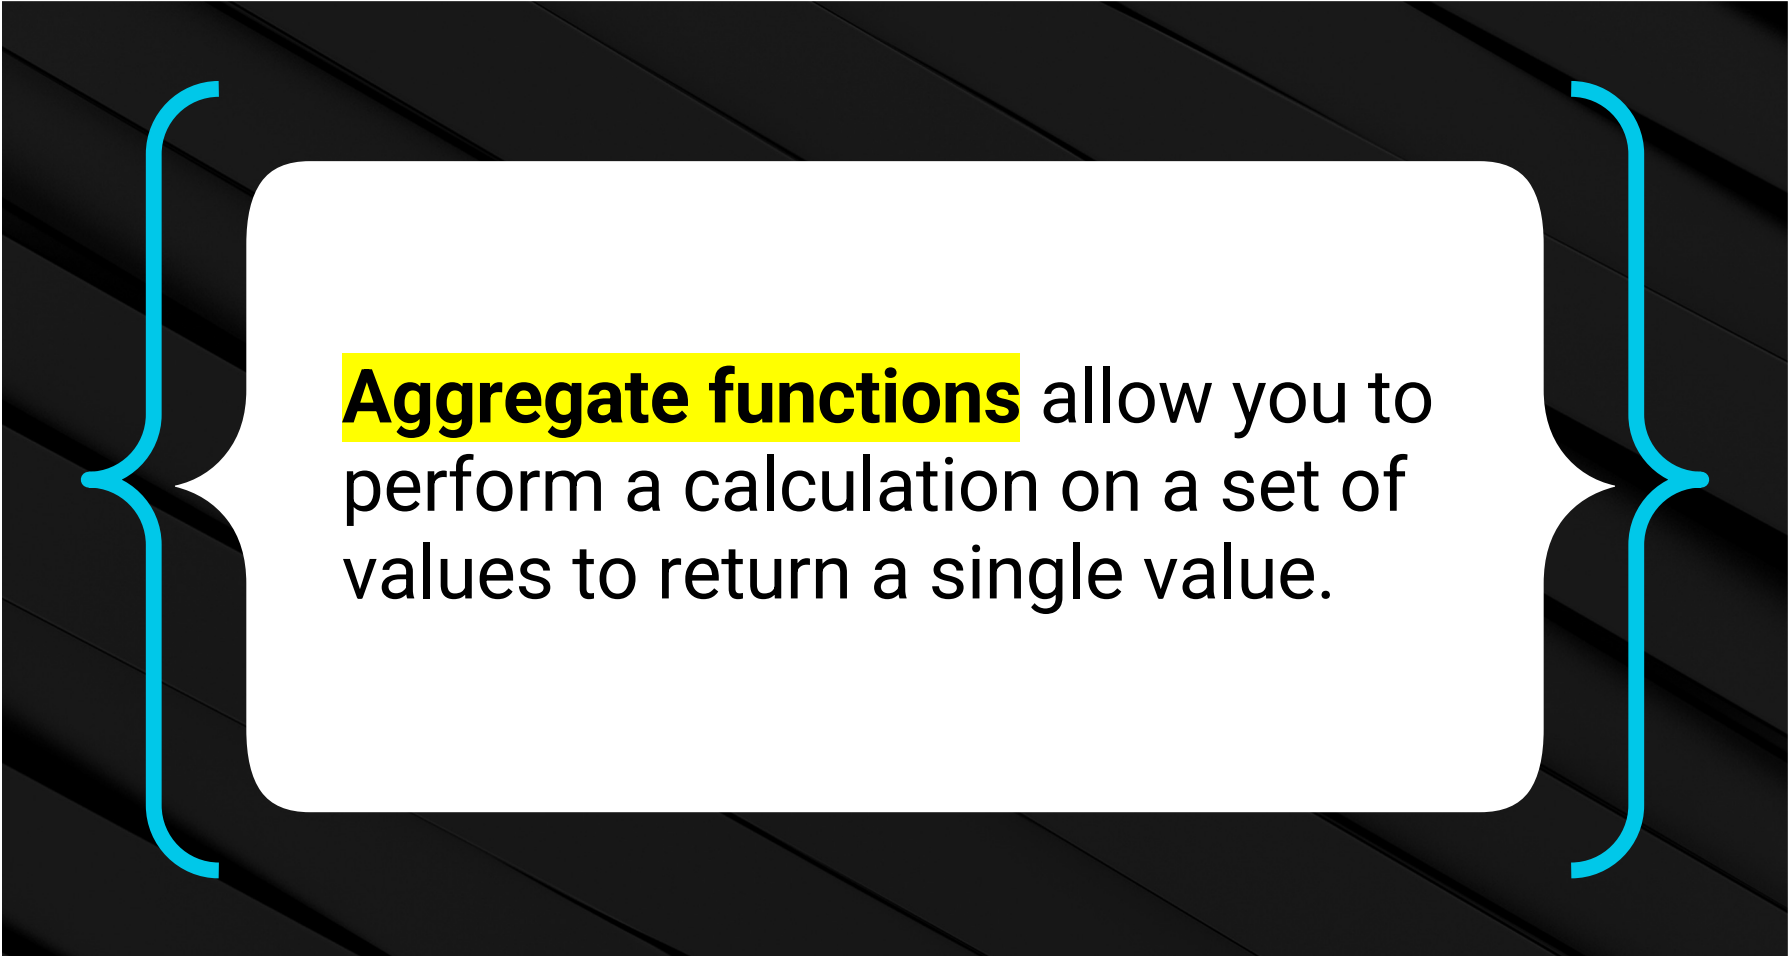
ord_no	purch_amt	ord_date	customer_id	salesman_id
70001	150.50	2012-10-05	3005	5002
70009	270.65	2012-09-10	3001	5005
70002	65.26	2012-10-05	3002	5001
70004	110.50	2012-08-17	3009	5003
70007	948.50	2012-09-10	3005	5002
70005	2400.60	2012-07-27	3007	5001
70008	5760.00	2012-09-10	3002	5001
70010	1983.43	2012-10-10	3004	5006
70003	2480.40	2012-10-10	3009	5003
70012	250.45	2012-06-27	3008	5002
70011	75.29	2012-08-17	3003	5007
70013	3045.60	2012-04-25	3002	5001

```
SELECT customer_id, MAX(purch_amt)
FROM orders
GROUP BY customer_id ;
```

customer_id	purch_amt
3005	150.5
3001	270.65
3002	65.26
3009	110.5
3005	948.5
3007	2400.6
3002	5760
3004	1983.43
3009	2480.4
3008	250.45
3003	75.29
3002	3045.6

customer_id	max
3004	1983.43
3008	250.45
3001	270.65
3007	2400.6
3005	948.5
3002	5760
3009	2480.4

Gregarious Aggregates



Aggregate functions allow you to perform a calculation on a set of values to return a single value.

Aggregate Functions

The most commonly used aggregate functions are:

AVG	Calculates the average of a set of values
COUNT	Counts the rows in a specific table or view
MIN	Returns the minimum value in a set of values
MAX	Returns the maximum value in a set of values
SUM	Calculates the sum of a set of values

Aggregate Functions

Aggregate functions are often used with:

01

The **GROUP BY** clause

02

The **HAVING** clause

03

The **SELECT** statement

Order By Aggregates

The **ORDER BY** function:



Is added towards the end of a query.



Returns in an ascending order by default.



Can also return in a descending order by adding **DESC**.



Can limit the return by adding **LIMIT**.



NOTE: Use the **ROUND** function to round up the number after the the decimal.



Pro Tip:

One important point about AGGREGATE FUNCTIONS:
The returned data is in random order!

Introduction to Subqueries

Subqueries

A subquery is nested inside a larger query. Subqueries occur in:

01

The **SELECT** statement

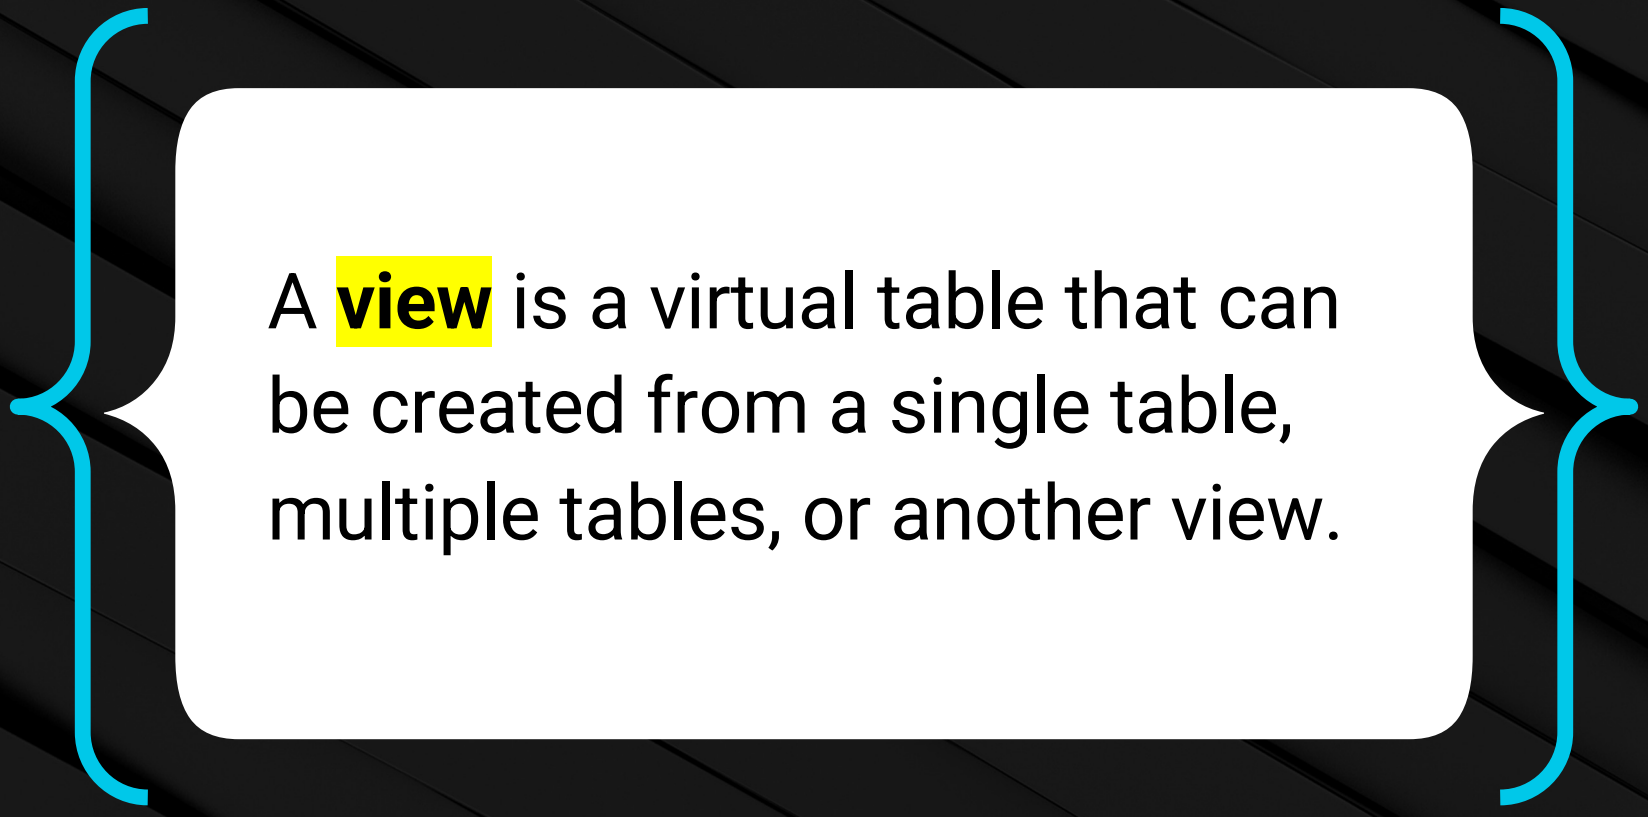
02

The **FROM** clause

03

The **WHERE** clause

SQL Views




A **view** is a virtual table that can be created from a single table, multiple tables, or another view.

SQL Views

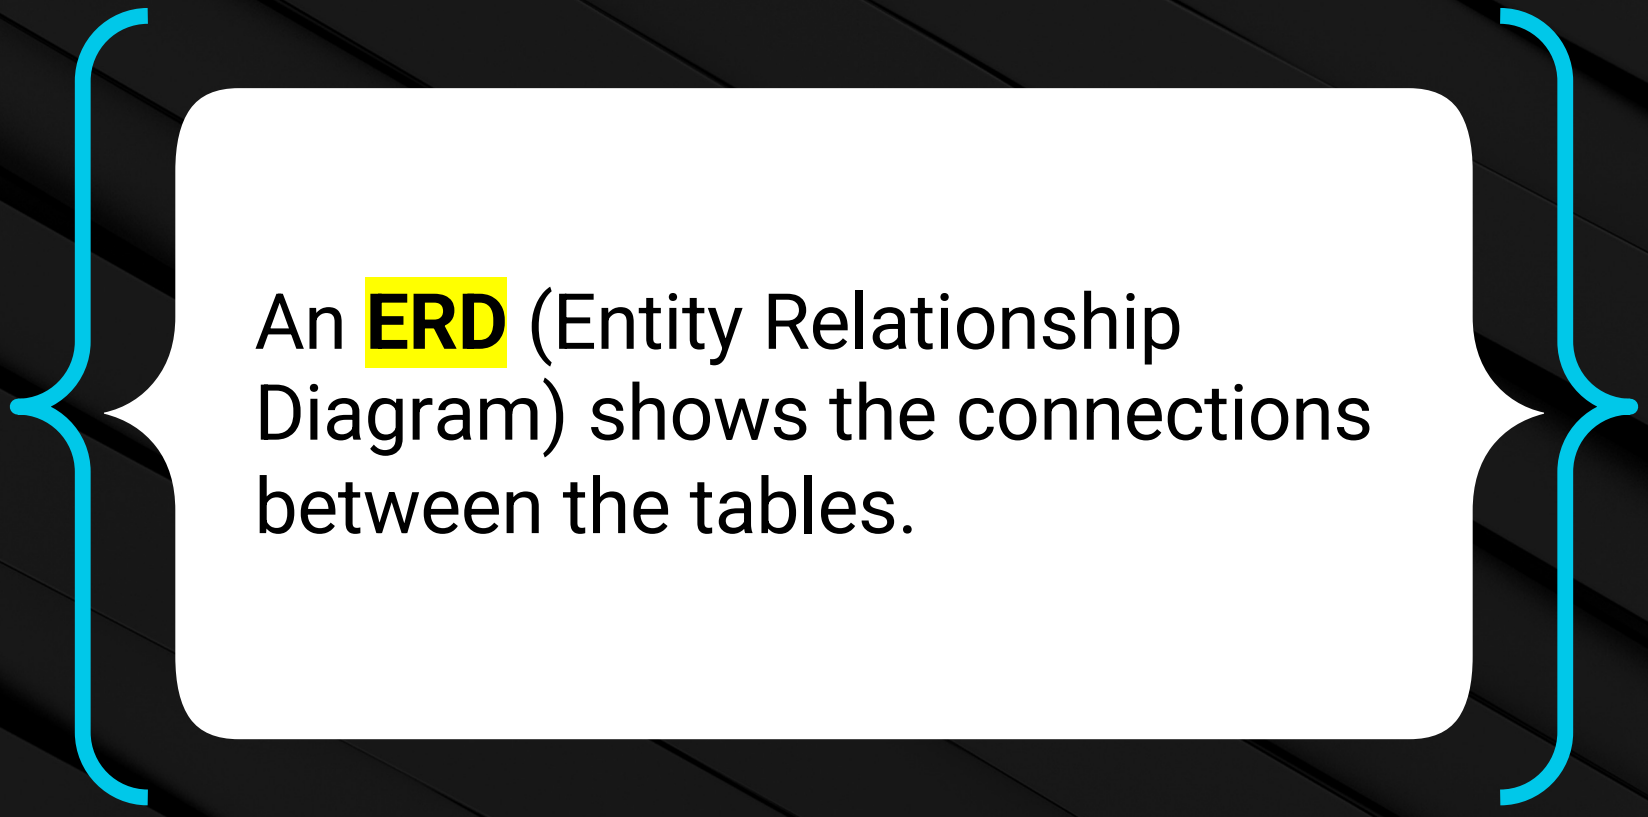
Views are created by using the `CREATE VIEW` statement.

Views are created from a single table, multiple tables, or another view.





How many people have rented
the film Blanket Beverly?



An **ERD** (Entity Relationship Diagram) shows the connections between the tables.

Entity Relationship Diagram


It's called an Entity Relationship Diagram because it shows:



The entities in your data model

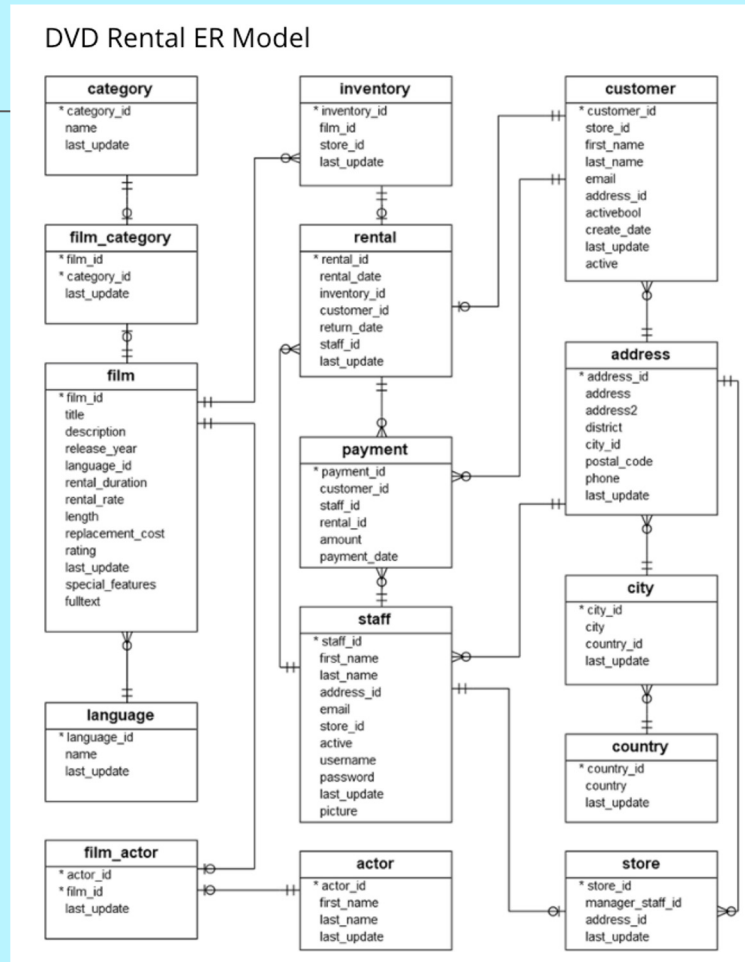


The relationship between entities



Entity Relationship Diagram

The schema makes it easier to identify the tables we need as well as the keys we will use to link our subqueries.



Foreign Keys

Foreign Keys

Foreign Keys reference the primary key of another table.

Can have a different name. It does not have to be unique.

Primary Key

	A	B
1	family_id	family
2	1	Smiths
3	2	Jones

Primary Key

Foreign Key

	A	B	C
1	child_id	family_id	children
2	11	1	Chris
3	22	1	Abby
4	33	1	Suzy

Data Relationships

Relationships Link Tables/Entities.

Types of relationships:

01

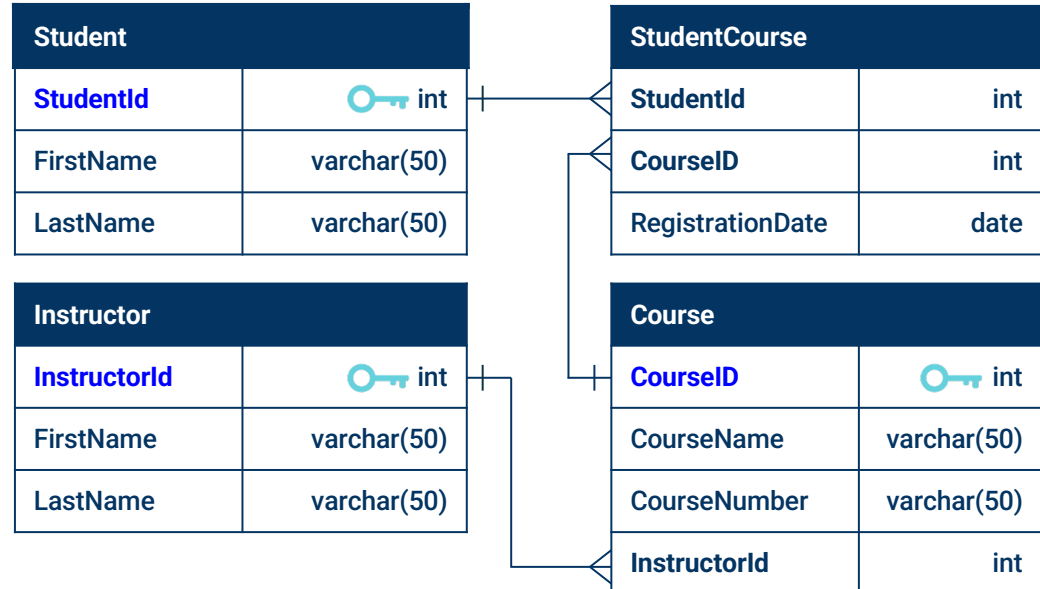
One-to-One

02

One-to-Many

03

Many-to-Many



One-to-One Relationship

Each item in one column is linked to only one other item from the other column.

ID	Name	Social Security
1	Homer	111111111
2	Marge	222222222
3	Lisa	333333333
4	Bart	444444444
5	Maggie	555555555

Here, each person in the Simpsons family can have only one social security number.

Each social security number can be assigned only to one person.

One-to-Many Relationship

This example has two tables. The first table lists only addresses.

The second table lists each person's Social Security number and address. As before, one Social Security number is unique to one individual.

ID	Address	ID	Name	Social Security	AddressID
11	742 Evergreen Terrace	1	Homer	111111111	11
12	221B Baker Street	2	Marge	222222222	11
		3	Lisa	333333333	11
		4	Bart	444444444	11
		5	Maggie	555555555	11
		6	Sherlock	112233445	12
		7	Watson	223344556	12

One-to-Many Relationship

- Each address can be associated with multiple people.
- Each person has an address.
- The two tables, joined, would look like this.

ID	Address	ID	Name	Social Security	AddressID
11	742 Evergreen Terrace	1	Homer	111111111	11
12	221B Baker Street	2	Marge	222222222	11
		3	Lisa	333333333	11
		4	Bart	444444444	11
		5	Maggie	555555555	11
		6	Sherlock	112233445	12
		7	Watson	223344556	12

Many-to-Many Relationship

- Each child can have more than one parent.
- Each parent can have more than one child.

ID	Child	ID	Parent
1	Bart	11	Homer
2	Lisa	12	Marge
3	Maggie		

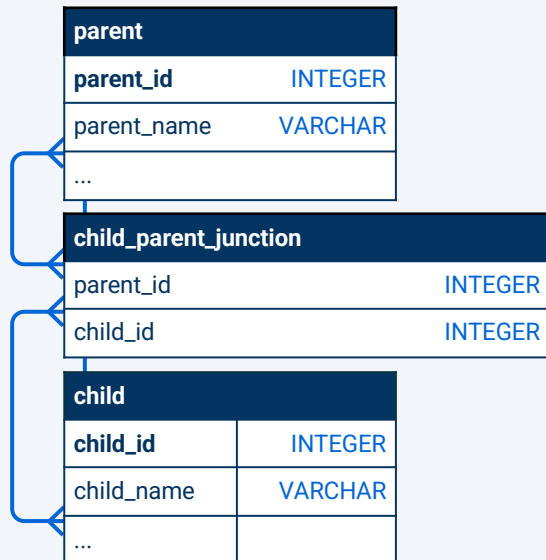
Many-to-Many Relationship

- Each child can have more than one parent.
- Each parent can have more than one child.
- The two tables are joined in a **junction table**.

ChildID	Child	ParentID	Parent
1	Bart	11	Homer
1	Bart	12	Marge
2	Lisa	11	Homer
2	Lisa	12	Marge
3	Maggie	11	Homer
3	Maggie	12	Marge

Junction Table

The junction table contains many parent_id's and many child_id's.



	parent_id integer	child_id integer
1	11	1
2	11	2
3	11	3
4	12	1
5	12	2
6	12	3

Join child and parent
table to junction table ↓

	parent_name character varying (255)	child_name character varying (255)
1	Homer	Bart
2	Homer	Lisa
3	Homer	Maggie
4	Marge	Bart
5	Marge	Lisa
6	Marge	Maggie

SQL Joins

INNER JOIN	Returns records that have matching values in both tables.
LEFT JOIN	Returns all records from the left table and the matched records from the right table.
RIGHT JOIN	Returns all records from the right table and the matched records from the left table.
CROSS JOIN	Returns records that match every row of the left table with every row of the right table. This type of join has the potential to make very large tables.
FULL OUTER JOIN	Places null values within the columns that do not match between the two tables, after an inner join is performed.